

Database Systems

Introduction to the course


Venkateswara Rao Kagita

Text Books

- **Adhsakkdi Y Raghuram Krishnan and Johannes Gehrke, *Database Management Systems*, 3/e, TMH, 2007.**
- Elamsri, Navathe, Somayajulu and Gupta, *Database Concepts*, Pearson Edition, 2006.
- Introduction to Database Systems, CJ Date, Pearson
- The Database Systems – The Complete Book, HG Molina, J D Ullman, J Widom Pearson
- Database Systems design, Implementation, and Management, Peter Rob & Carlos Coronel 7th Edition.
-

Class Rules

- Assignments are to be completed individually
- Academic **honesty** taken seriously



Syllabus

Topic	#hours
• Overview of Database Systems	4
• Introduction to Database Design, ER-Model	4
• Relational model	4
• Relational algebra and Calculus	5
• SQL	6
• Functional Dependencies and normalization process	4
• Transaction Management	5
• Concurrency Control	4
• Crash recovery	3
• Security and Authorization	3
Total	42

Overview of Database Systems

What is Database System ?

- **Data:**
 - Known facts that can be recorded and have an implicit meaning.
- **Database:**
 - A collection of related data.
- **Database Management System (DBMS):**
 - Software for creating, storing, maintaining, and accessing database files
 - Makes using databases more efficient
 - There are many databases available like MySQL, Sybase, Oracle, MongoDB, Informix, PostgreSQL, SQL Server, etc.
- **Database System:**
 - The DBMS software together with the data itself. Sometimes, the applications are also included.

What is the need of DBMS?

- When dealing with huge amount of data, there are two things that require optimization:
 - **Data storage:** Efficient storage, no redundancy.
 - **Data retrieval:** Retrieve the data quickly when needed
- Purpose of Database Systems
 - The main purpose of database systems is to manage the data.
 - To manage the data we need to store this data somewhere where we can add new data, delete unused data, update outdated data, retrieve data.
 - To perform these operations on data we need a Database management system that allows us to store the data in such a way so that all these operations can be performed on the data efficiently.
- Why Use a DBMS ?
 - Large amount of data (Giga's, Tera's) drives many business processes
 - Data is very structured
 - Persistent and valuable data
 - Performance requirements
 - Concurrent and restricted access to the data

Simplified database system environment

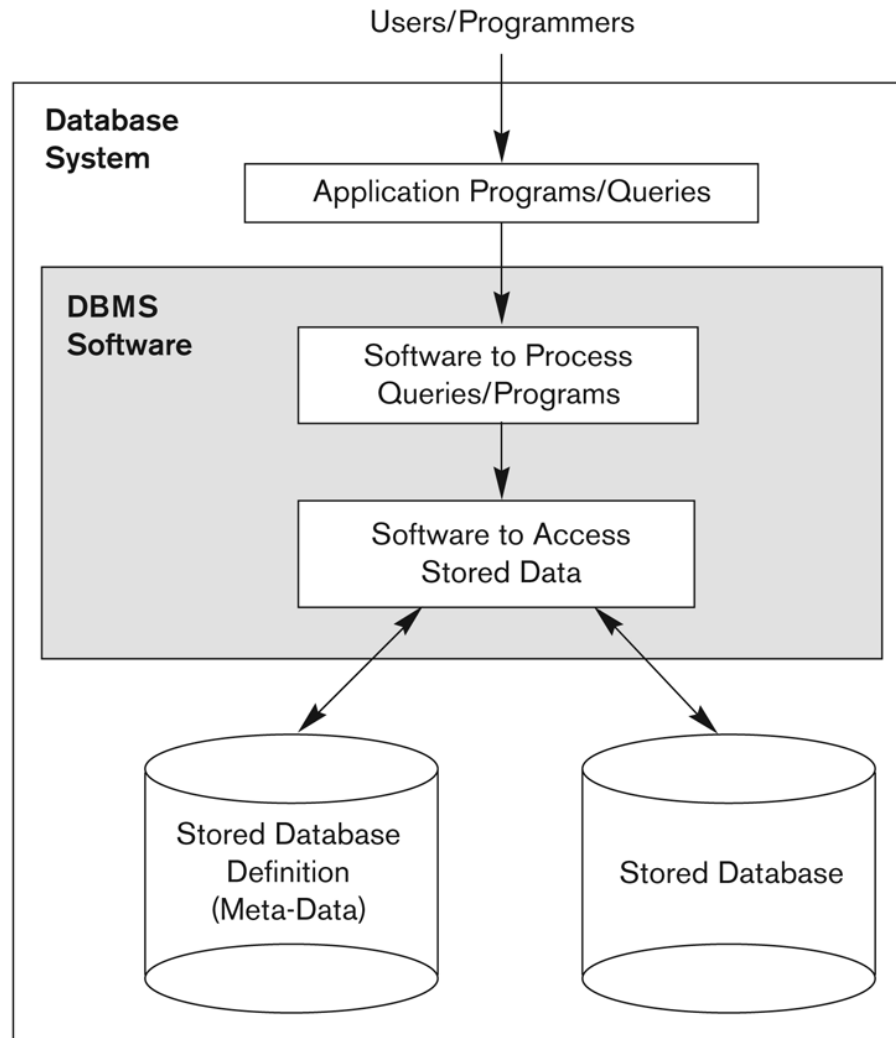
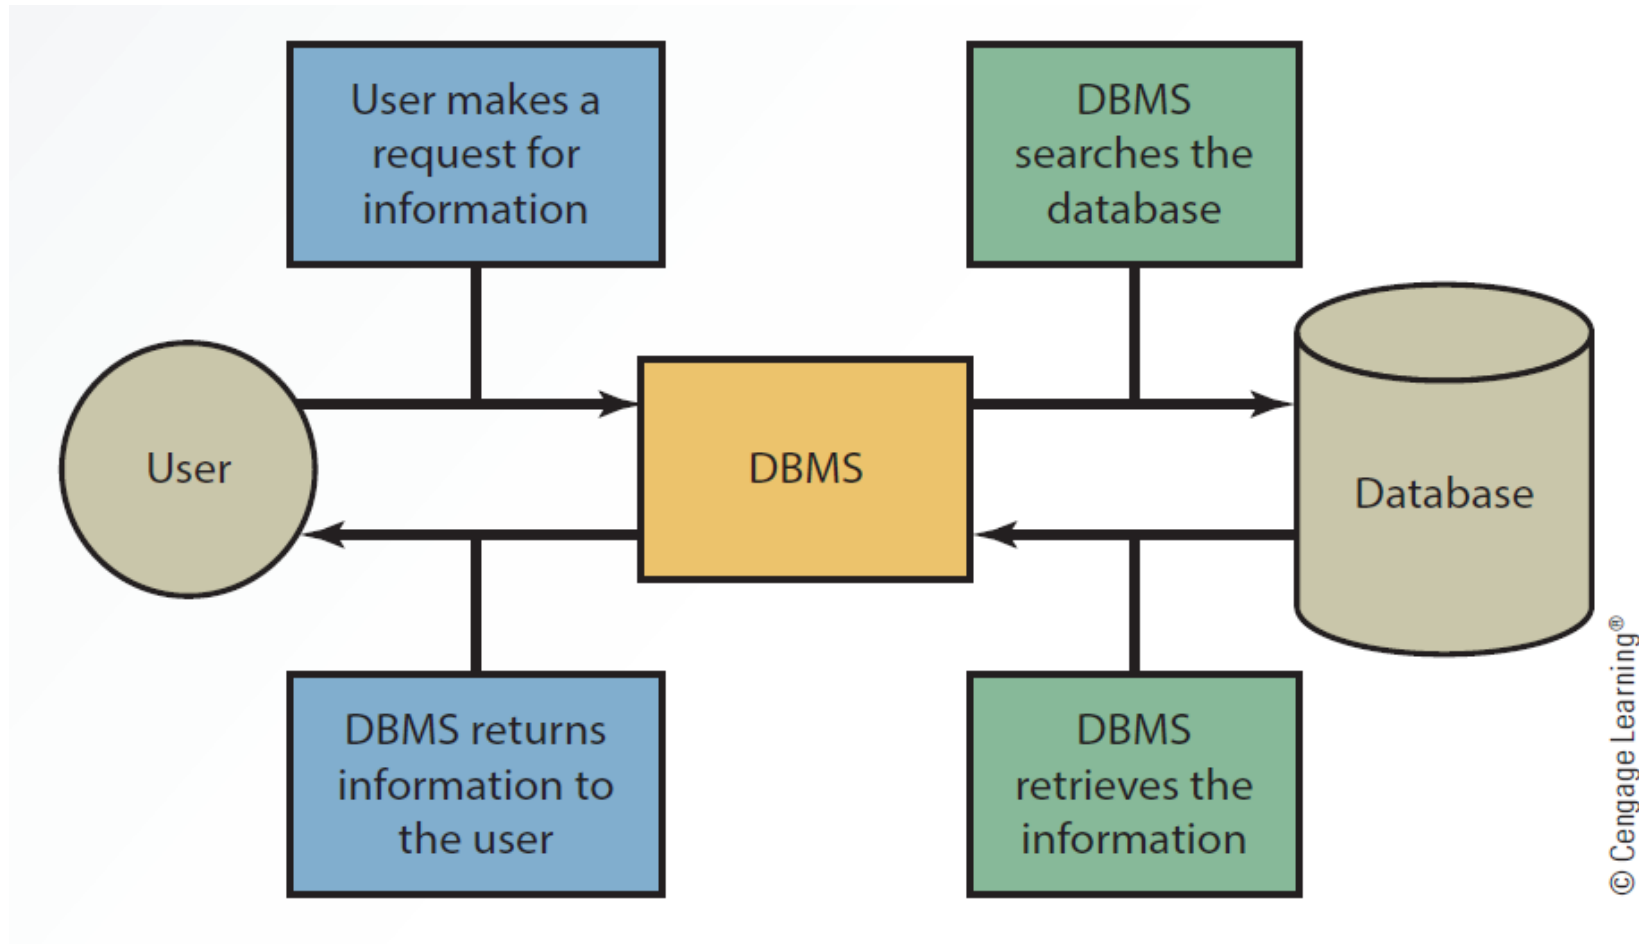


Figure 1.1
A simplified database
system environment.


Interaction Between the User, DBMS and Database



A few of Applications

- **Telecom:** Calls made, network usage, customer details etc.
- **Industry:** Where it is a manufacturing unit, to keep the records of ins and outs.
- **Banking System:** For storing customer info, tracking day to day credit and debit transactions, generating bank statements etc.
- **Sales:** To store customer information, production information and invoice details.
- **Airlines:** Reservation information along with flight schedule is stored in database.
- **Education sector:** To store and retrieve the data regarding student details, staff details, course details, exam details, payroll data, attendance details, fees details etc.
- **Online shopping:** To store the product information, your addresses and preferences, credit details and provide you the relevant list of products based on your query.

Types of Database



Why not use simple file to
store the data ?

Drawbacks of File system

- **Data redundancy**
 - Duplication of data
 - Requires more storage than needed.
 - Higher storage costs and poor access time.
- **Data inconsistency:**
 - Data redundancy leads to data inconsistency
- **Data Isolation:** a property that determines when and how changes made by one operation become visible to other concurrent users and systems. Difficult to achieve in file processing system.
- **Dependency on application programs:** Changing files would lead to change in application programs.
- **Atomicity issues:** “All or nothing”
 - For example: Ram transfers 1000 INR to Krishna account.
 - Debt 1000 rupees from Ram account
 - Credit 1000 rupees to Krishna account
 - What if the system fails after first operation ?
 - In such case the rollback of operation should occur to maintain the atomicity of transaction.
 - It is **difficult to achieve atomicity in file processing systems.**
- **Data Security:**
 - Data should be secured from unauthorized access.
 - For example a student should not be able to access grades.
 - Such kind of security constraints are difficult to apply in file processing systems.

Advantage of DBMS over file system

- **Data Independence**
 - Application programs should not, ideally, be exposed to details of data representation and storage.
- **No redundant data:**
 - Data normalization removes the data redundancy.
 - Saves storage and improves access time.
- **Data Consistency**
- **Data Integrity and Security**
 - data is always accessed through the DBMS, the DBMS can enforce integrity constraints.
 - Each user has a different set of access thus data is secured.
- **Privacy:** Limited access means privacy of data.
- **Easy access to data** – Data is easily accessible with fast response times.
- **Concurrent Access**
- **Easy recovery:**
 - Keeps the backup of data.
 - Easy to recover data in case of a failure.
- **Flexible:** Database systems are more flexible than file processing systems.


Disadvantages of DBMS:

- DBMS implementation cost is high compared to the file system
- Complexity: Database systems are complex to understand
- Performance: Database systems are generic, making them suitable for various applications. However this feature affect their performance for some applications.

When not to use a DBMS

- Main inhibitors (costs) of using a DBMS:
 - High initial investment and possible need for additional hardware.
 - Overhead for providing generality, security, concurrency control, recovery, and integrity functions.
- When a DBMS may be unnecessary:
 - If the database and applications are simple, well defined, and not expected to change.
 - If there are stringent real-time requirements that may not be met because of DBMS overhead.
 - If access to data by multiple users is not required.
- When no DBMS may suffice:
 - If the database system is not able to handle the complexity of data because of modeling limitations
 - If the database users need special operations not supported by the DBMS.

DBMS components

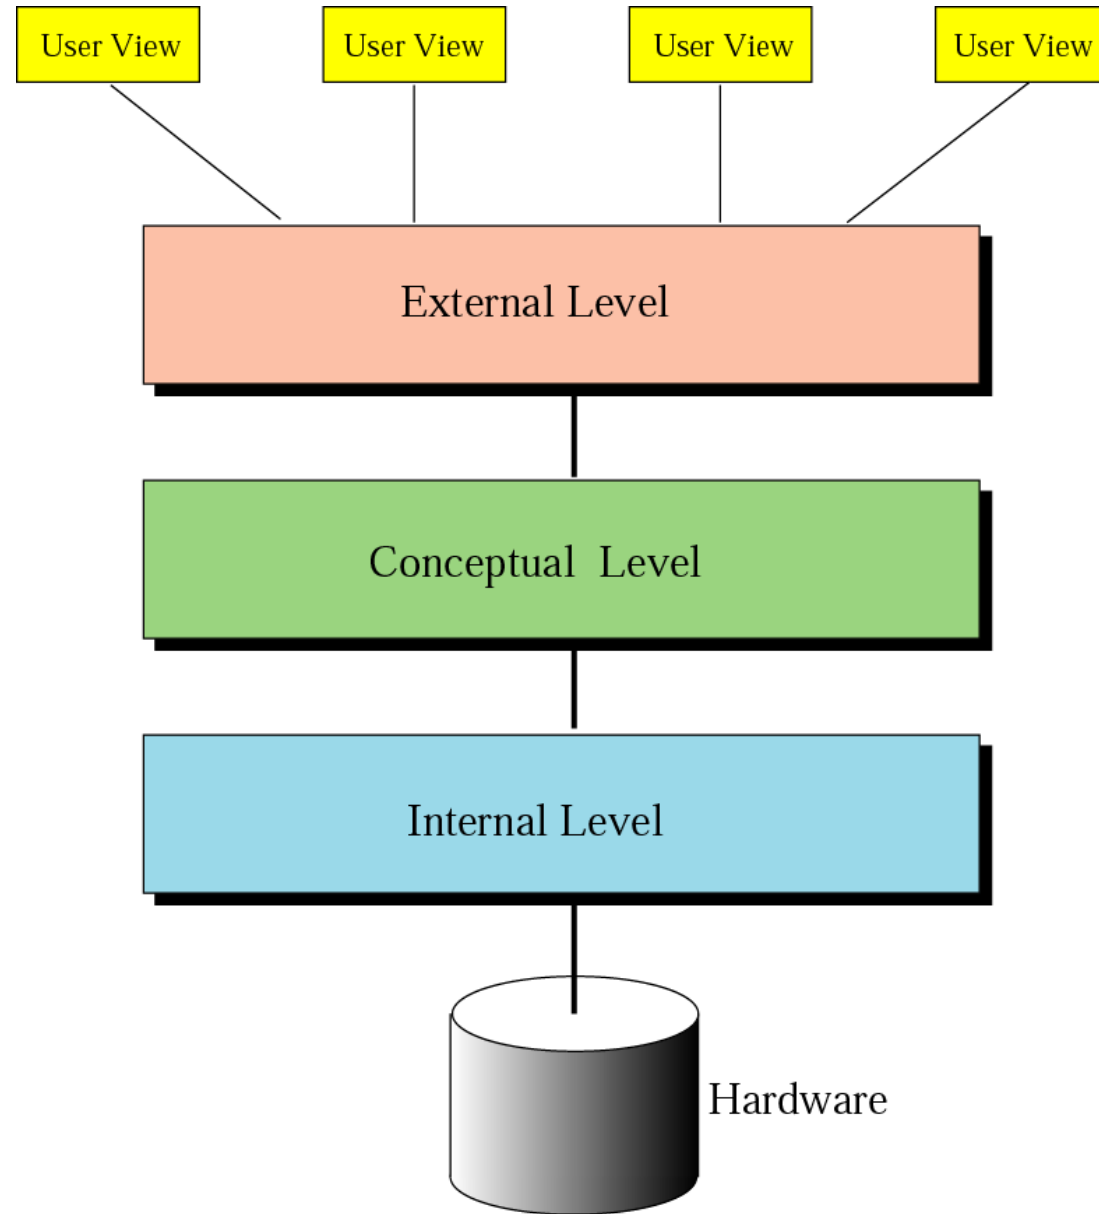


- **Hardware** –
the physical computer system that allows physical access to data.
- **Software** –
the actual program that allows users to access, maintain, and update physical data.
- **Data** – stored physically on the storage devices
- **Users** –
 - database implementors, who build DBMS software,
 - End users - Normal user and DBA (Database Administrator)
 - Application programs, Application programmers
- **Procedures** –
a set of rules that should be clearly defined and followed by the users.

Database Administrator Responsibilities

- Design of the Conceptual and Physical Schemas
- Security and Authorization
- Data Availability and Recovery from Failures.
- Database Tuning.

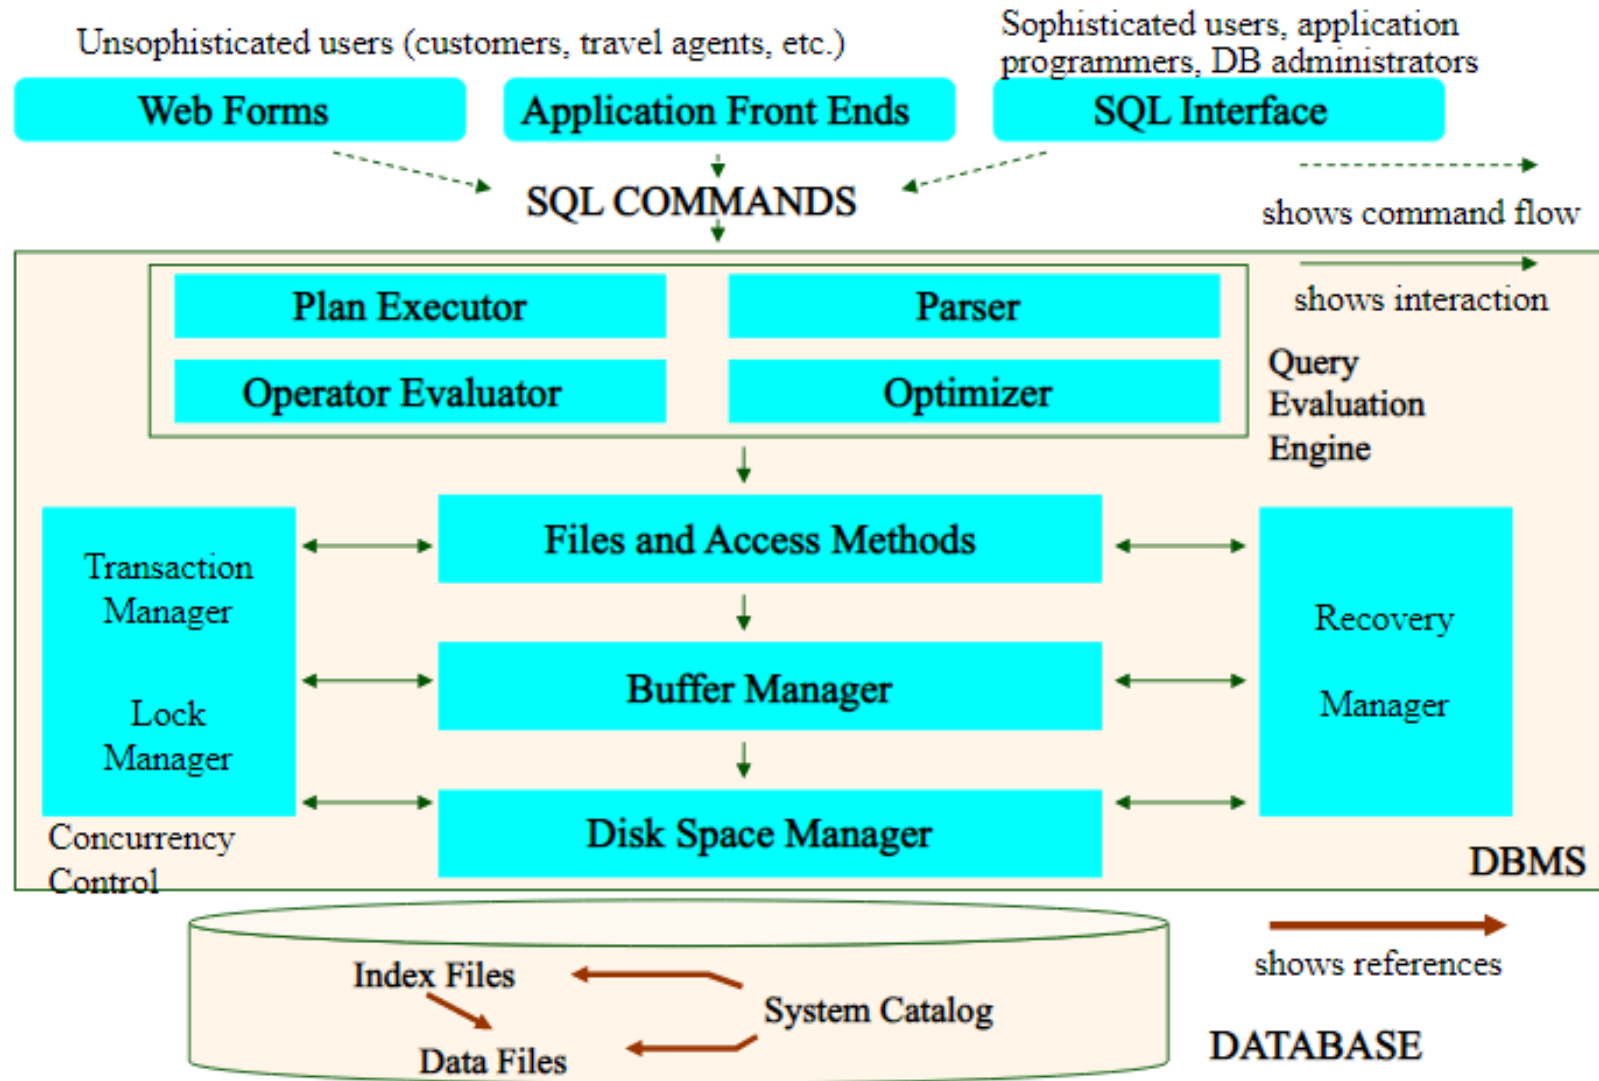
Database Architecture



Architecture

- Internal level –
 - Determines where data are actually stored on the storage device.
 - Low-level access method
- Conceptual level –
 - Defines the logical view of the data
 - The main functions of DBMS are in this level.
- External level –
 - Interacts directly with the user.
 - Change the data coming from the conceptual level to a format and view that are familiar to the users.

DBMS Architecture: Revisited

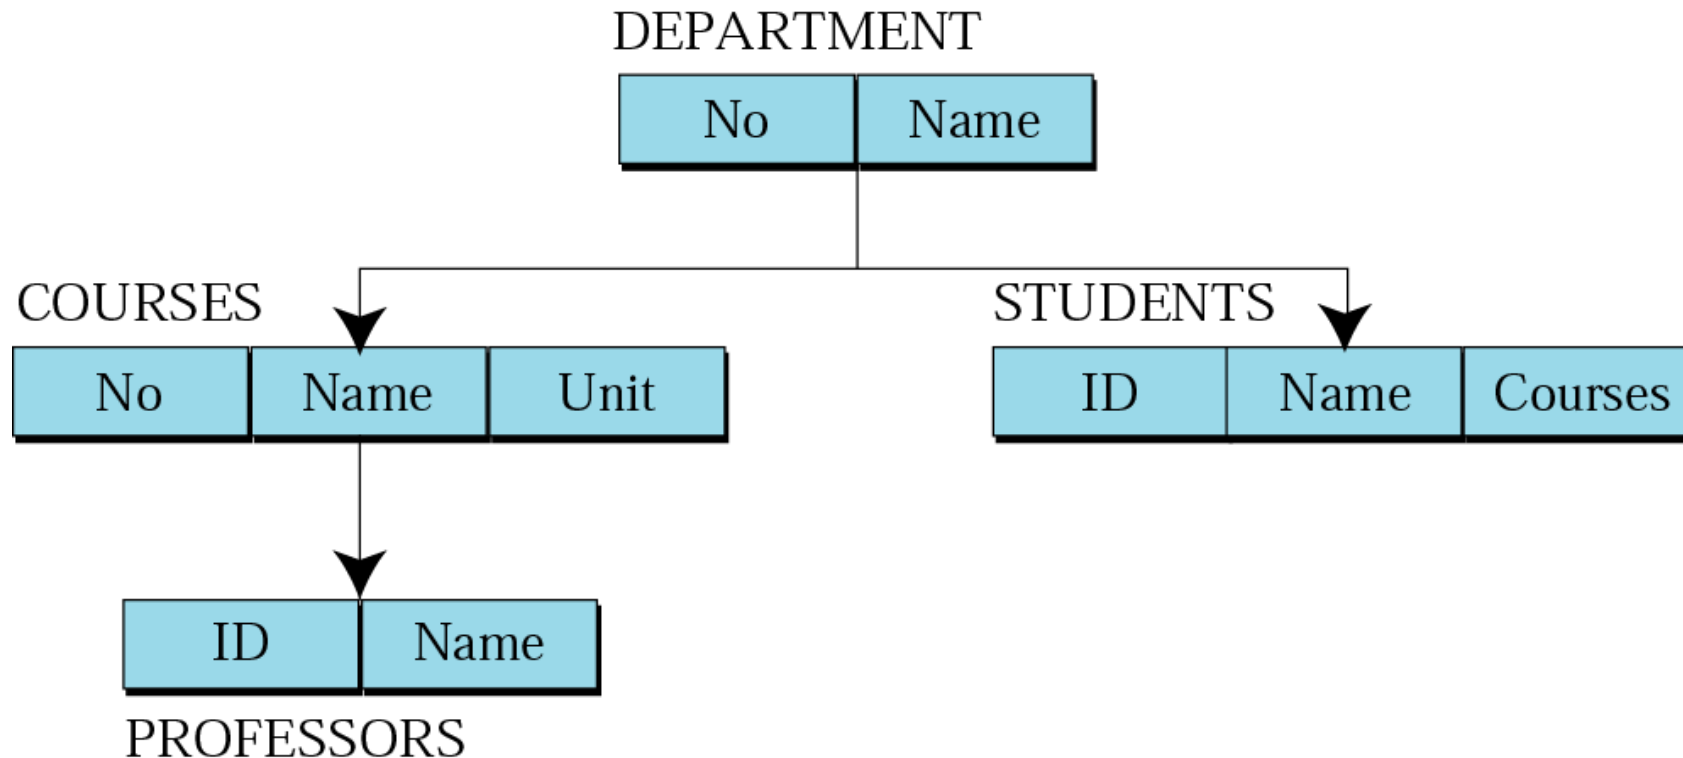


Database models

- A **data model** is a collection of high-level data description constructs that hide many low level storage details
- Defines the logical design of data.
- Describes the relationships between different parts of data.
- Different models
 - Hierarchical model (e.g., used in IBM's IMS DBMS)
 - Network model (e.g., used in IDS and IDMS)
 - Relational model (e.g. Oracle, Sybase, Microsoft's Access, etc.)
 - The object-oriented model (e.g., Objectstore and Versant)
 - The object-relational model (e.g., used in DBMS products from IBM, Informix, ObjectStore, Oracle, Versant, and others).
 - ...
- While many databases use the hierarchical and network models and systems based on the object-oriented and object-relational models are gaining acceptance in the marketplace, the dominant model is the relational model.

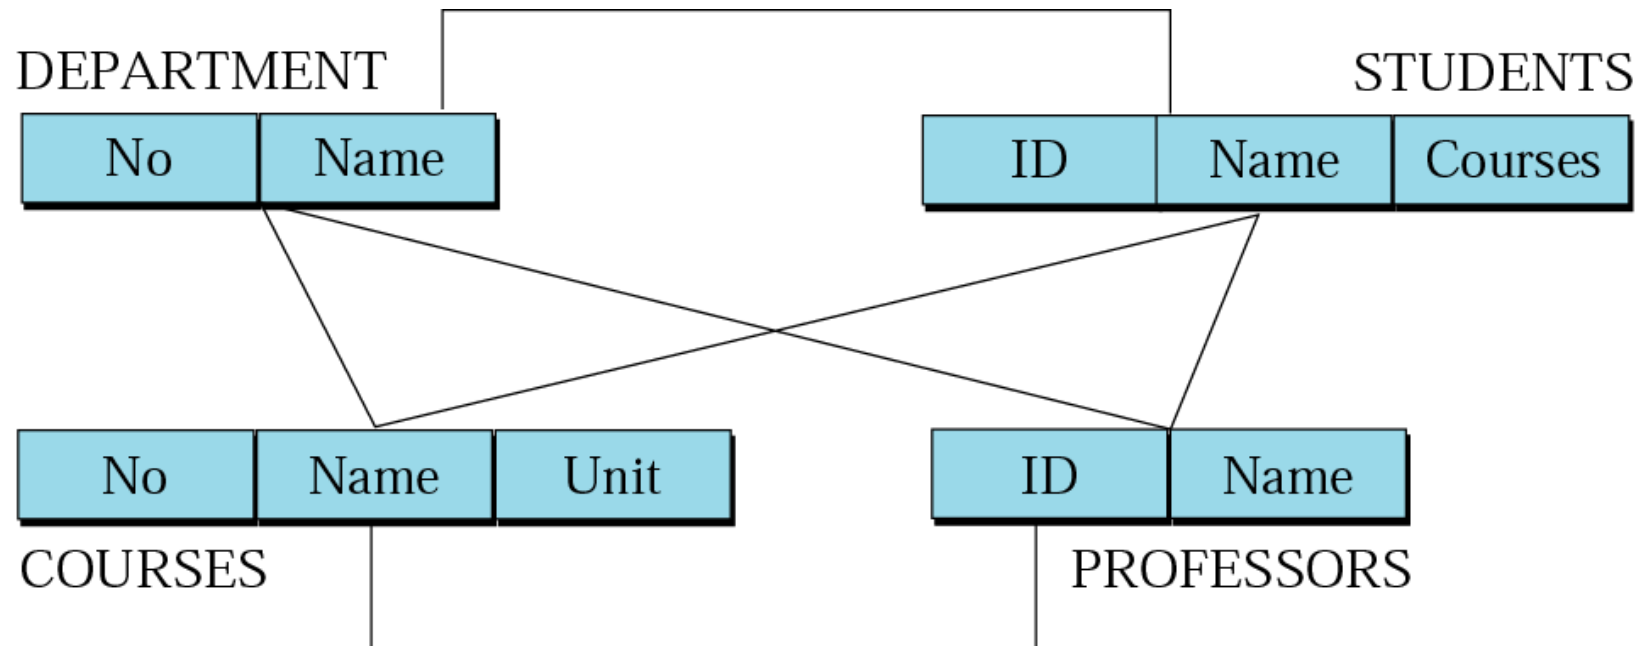
Hierarchical model

- Data are organized as an upside down tree.
- Each entity has only one parent but can have several children.



Network model

- The entities are organized in a **graph**.
- Some entities can be accessed through several paths.



Relational model

- Data are organized in two-dimensional tables called **relations**.
- The tables are related to each other.
- The most popular model. We mainly focus on relational model in this course

DEPARTMENT

No	Name
...	...
...	...
...	...

PROFESSORS

ID	Name	Dept-No	Courses
...
...
...
...

COURSES

No	Dept-No	Prof-ID	Unit
...
...
...
...

STUDENTS

ID	Name	Courses
...
...
...
...

Relational model

- RDBMS (Relational Database Management System)
- external view
 - The data are represented as a set of relations.
 - A **relation** is a two-dimensional table.
- This **doesn't** mean that data are stored as tables; the physical storage of the data is **independent** of the way the data are logically organized.

Relation

- **Name** – each relation in a relational database should have a name that is **unique** among other relations.
- **Attribute** – each **column** in a relation.
 - The degree of the relation – the total number of attributes for a relation.
- **Tuple** – each **row** in a relation.
 - The cardinality of the relation – the total number of rows in a relation.

Example:

Fields (Attributes, Columns)

sid	Name	login	age	Gross
1111	Dave	dave@cs	19	1.2
2222	Jones	Jones@cs	18	2.3
333	Smith	smith@ee	18	3.4
4444	Smith	smith@math	19	4.5

Field names

Tuples (Records, Rows)

Students

Operations on relations

- In a relational database, we can define several operations to create new relations out of the existing ones.
- Basic operations:
 - Insert
 - Delete
 - Update
 - Select
 - Project
 - Join
 - Union
 - Intersection
 - Difference

Insert operation

- An **unary** operation.
- Insert a new **tuple** into the relation.

COURSES

No	Course-Name	Unit
CIS15	Intro to C	5
CIS17	Intro to Java	5
CIS19	UNIX	4
CIS51	Networking	5

Insert


No	Course-Name	Unit
CIS15	Intro to C	5
CIS17	Intro to Java	5
CIS19	UNIX	4
CIS51	Networking	5
<i>CIS52</i>	<i>TCP/IP Protocols</i>	<i>6</i>

Delete operation

- An **unary** operation.
- Delete a **tuple** defined by a **criterion** from the relation.

COURSES

No	Course-Name	Unit
CIS15	Intro to C	5
CIS17	Intro to Java	5
CIS19	UNIX	4
CIS51	Networking	5
CIS52	TCP/IP Protocols	6




No	Course-Name	Unit
CIS15	Intro to C	5
CIS17	Intro to Java	5
CIS51	Networking	5
CIS52	TCP/IP Protocols	6

Update operation

- An **unary** operation.
- Changes the value of some attributes of a tuple.

COURSES



Select operation

- An **unary** operation.
- It is applied to one single relation and creates another relation.
- The tuples in the resulting relation are a subset of the tuples in the original relation.
- Use some **criteria** to select

COURSES

No	Course-Name	Unit
CIS15	Intro to C	5
CIS17	Intro to Java	5
CIS19	UNIX	4
CIS51	Networking	5
CIS52	TCP/IP Protocols	6

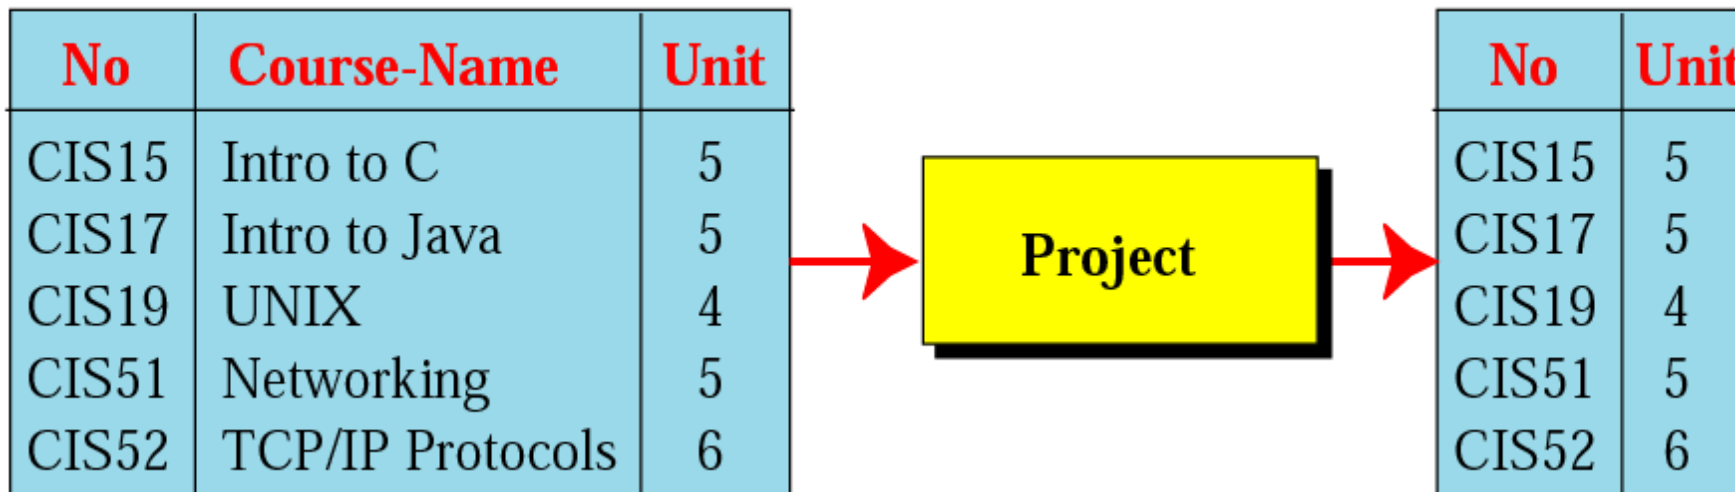


No	Course-Name	Unit
CIS15	Intro to C	5
CIS17	Intro to Java	5
CIS51	Networking	5

Project operation

- An **unary** operation.
- It is applied to one single relation and creates another relation.
- The attributes in the resulting relation are a subset of the attributes in the original relation.

COURSES



Join operation

- A **binary** operation.
- Combines two relations based on common attributes.

COURSES

No	Course-Name	Unit
CIS15	Intro to C	5
CIS17	Intro to Java	5
CIS19	UNIX	4
CIS51	Networking	5
CIS52	TCP/IP Protocols	6

TAUGHT-BY

No	Professor
CIS15	Lee
CIS17	Lu
CIS19	Walter
CIS51	Lu
CIS52	Lee

Join

No	Course-Name	Unit	Professor
CIS15	Intro to C	5	Lee
CIS17	Intro to Java	5	Lu
CIS19	UNIX	4	Walter
CIS51	Networking	5	Lu
CIS52	TCP/IP Protocols	6	Lee

Union operation

- A **binary** operation.
- Creates a new relation in which each tuple is either in the first relation, in the second, or in both.
- The two relations must have the same attributes.

CIS15-Roster

Student-ID	F-Name	L-Name
145-67-6754	John	Brown
232-56-5690	George	Yellow
345-89-6580	Anne	Green
459-98-6789	Ted	Purple

CIS52-Roster

Student-ID	F-Name	L-Name
342-88-9999	Rich	White
145-67-6754	John	Brown
232-56-5690	George	Yellow

Union

Student-ID	F-Name	L-Name
145-67-6754	John	Brown
232-56-5690	George	Yellow
345-89-6580	Anne	Green
459-98-6789	Ted	Purple
342-88-9999	Rich	White

Intersection operation

- A **binary** operation.
- Creates a new relation in which each tuple is a member in both relations.
- The two relations must have the same attributes.

CIS15-Roster

Student-ID	F-Name	L-Name
145-67-6754	John	Brown
232-56-5690	George	Yellow
345-89-6580	Anne	Green
459-98-6789	Ted	Purple

CIS52-Roster

Student-ID	F-Name	L-Name
342-88-9999	Rich	White
145-67-6754	John	Brown
232-56-5690	George	Yellow

Intersection

Student-ID	F-Name	L-Name
145-67-6754	John	Brown
232-56-5690	George	Yellow

Difference operation

- A **binary** operation.
- Creates a new relation in which each tuple is in the first relation but not the second.
- The two relations must have the same attributes.

CIS15-Roster

Student-ID	F-Name	L-Name
145-67-6754	John	Brown
232-56-5690	George	Yellow
345-89-6580	Anne	Green
459-98-6789	Ted	Purple

CIS52-Roster

Student-ID	F-Name	L-Name
342-88-9999	Rich	White
145-67-6754	John	Brown
232-56-5690	George	Yellow

Difference

Student-ID	F-Name	L-Name
345-89-6580	Anne	Green
459-98-6789	Ted	Purple

Schema

- A description of data in terms of a data model is called a schema.
- In the relational model, the schema for a relation specifies
 - Relation name
 - The name of each field (or attribute or column)
 - The type of each field.
- Example: Student information in a university
 - Schema: Students(*sid*: string, *name*: string, *login*: string, *age*: integer, *gpa*: real)

Example:

Fields (Attributes, Columns)

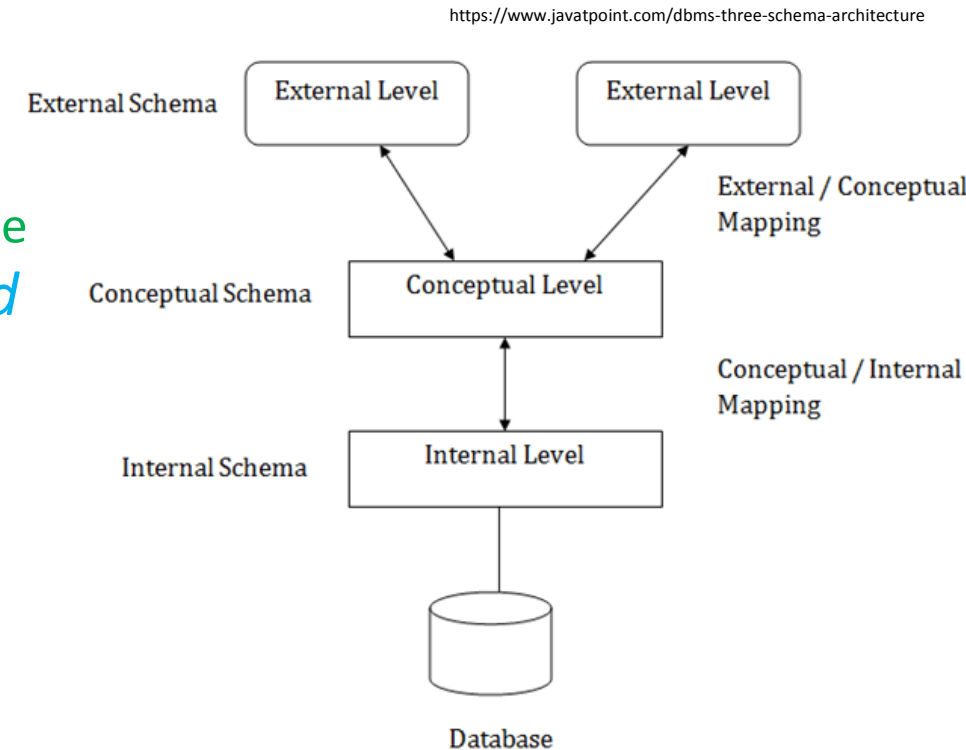
sid	Name	login	age	Gross
1111	Dave	dave@cs	19	1.2
2222	Jones	Jones@cs	18	2.3
333	Smith	smith@ee	18	3.4
4444	Smith	smith@math	19	4.5

Field names

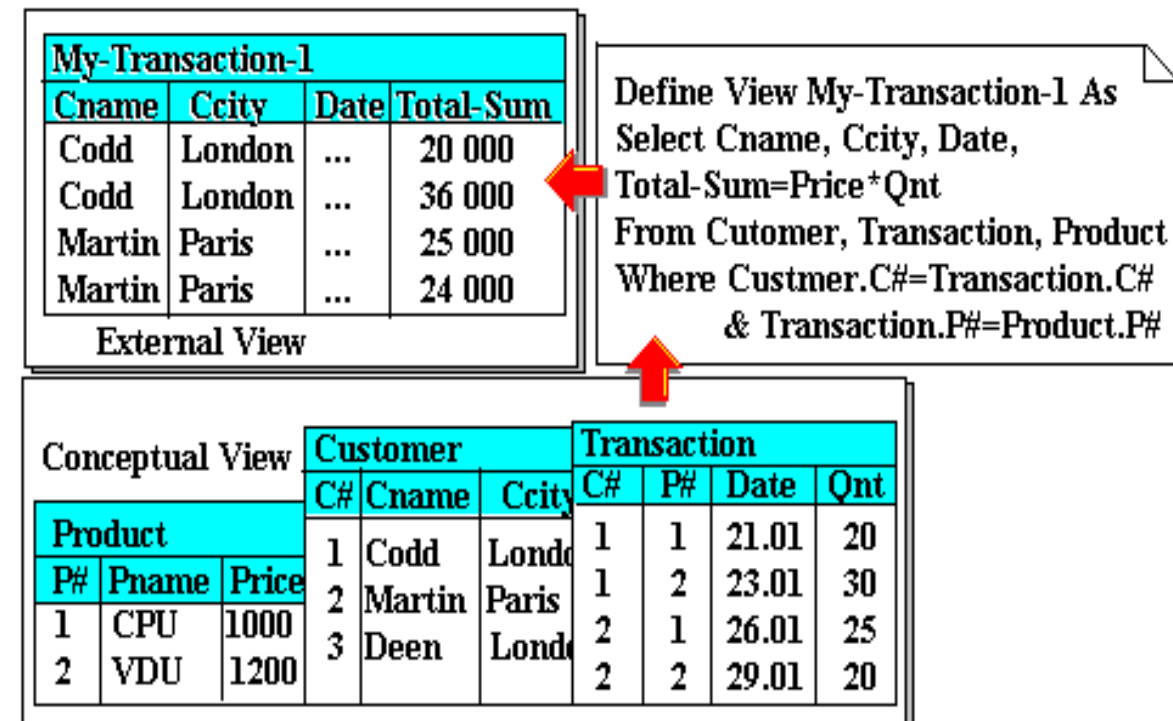
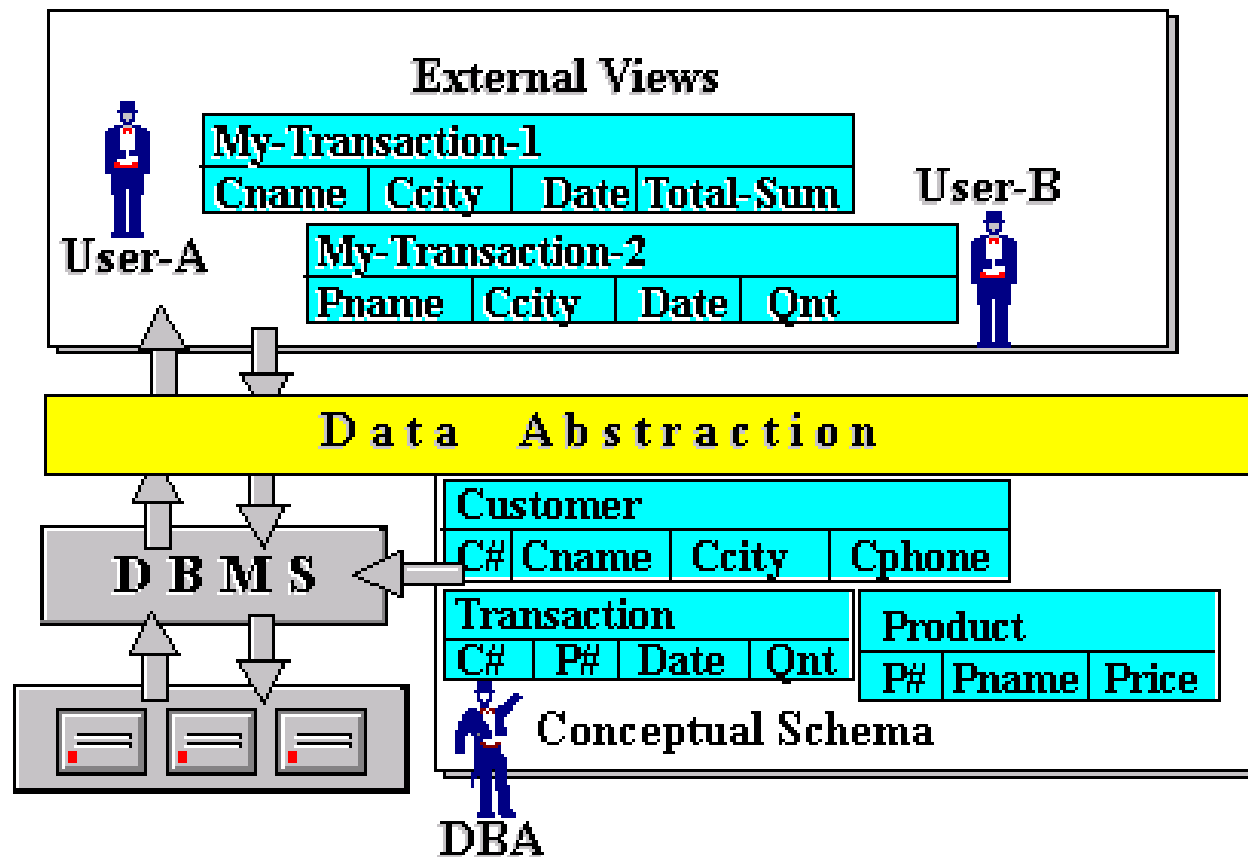
Tuples (Records, Rows)

Levels of Abstraction in a DBMS

- The data in a DBMS is described at three levels of abstraction
 - *External schema: How users view the data*
 - *Eg: View number of enrollements in each course*
Courseinfo(cid: string, cname: string, enrollment: integer)
 - *Conceptual schema: Logical structure, tables and their attributes*
 - *Eg: Students (sid: string, name: string, ..., gpa:real)*
Courses (cid: string, cname: string, credits: integer)
Enrolled(sid: string, cid: string, grade: string)
 - *Physical schema: Storage methods and access methods*
 - *Relations stored as unordered files*
 - *Index on first column of students*
- A data definition language (DDL) is used to define the external and conceptual schemas.



Example



Data Independence

- Change the Database schema at one level of a database system without requiring to change the schema at the next higher level.
 - **Physical Data Independence:** With Physical independence, you can easily change the physical storage structures or devices with no effect on the conceptual schema.
 - Any change done would be absorbed by the mapping between the conceptual and internal levels
 - Examples of internal level changes: Due to Physical independence, any these changes will not affect the conceptual layer.
 - Using a new storage device like Hard Drive or Magnetic Tapes
 - Modifying the file organization technique in the Database
 - Switching to different data structures.
 - Changing the access method, Modifying indexes.
 - Changes to compression techniques or hashing algorithms.
 - Change of Location of Database from say C drive to D Drive
 - **Logical Data Independence:** Logical Data Independence is the ability to change the conceptual scheme without changing External views or External programs.
 - Any change made will be absorbed by the mapping between external and conceptual level.
 - Examples of Conceptual level changes: Due to Logical independence, any of these changes will not affect the external layer.
 - Add/Modify/Delete a new attribute, entity or relationship is possible without a rewrite of existing application programs
 - Merging two records into one, - Breaking an existing record into two or more records

Transaction Management

- Scenario 1: Multiple customers accessing the available seat information in Airlines and trying to book.

- Scenario 2: Two transactions

T_1 :
 Read(x)
 $x \leftarrow x+100$
 Write(x)
 Commit

T_2 :
 Read(x)
 $x \leftarrow x-100$
 Write(x)
 Commit

Possible execution sequences:

T_1 : Read(x)
 T_1 : $x \leftarrow x+100$
 T_1 : Write(x)
 T_1 : Commit
 T_2 : Read(x)
 T_2 : $x \leftarrow x-100$
 T_2 : Write(x)
 T_2 : Commit

T_1 : Read(x) 500
 T_2 : Read(x) 500
 T_1 : $x \leftarrow x+100$ 600
 T_1 : Write(x) 600
 T_2 : $x \leftarrow x-100$ 400
 T_2 : Write(x) 400
 T_1 : Commit 600
 T_2 : Commit 400

Principles of Transaction

ATOMICITY

- all or nothing

CONSISTENCY

- no violation of integrity constraints

ISOLATION

- concurrent changes invisible \Rightarrow serializable

DURABILITY

- committed updates persist

Before: X : 500	Y: 200
Transaction T	
T1	T2
Read (X) X: = X - 100 Write (X)	Read (Y) Y: = Y + 100 Write (Y)
After: X : 400	Y : 300

T	T''
Read (X) X: = X*100 Write (X) Read (Y) Y: = Y - 50 Write	Read (X) Read (Y) Z: = X + Y Write (Z)

Queries in a DBMS

- Here are some questions a user might ask from a university database system.
 - What is the name of the student with student ID 1234567
 - What is the average salary of professors who teach course CS5647
 - How many students are enrolled in CS5647
 - What fraction of students in CS564 received a grade better than B7
 - Is any student with a CPA less than 3.0 enrolled in CS5647
- Such questions involving the data stored in a DBMS are called queries.
 - A DBMS provides a specialized language, called the query language, in which queries can be posed. Eg: SQL (structured query language)
- A DBMS enables users to create, modify, and query data through a data manipulation language (DML).

References

- Adhsakkdi Y Raghuram Krishnan and Johannes Gehrke, Database Management Systems, 3/e, TMH, 2007.
- Referred various online sources for examples and images.